

dSprites VSA

1 Introduction

The goal of this work is to extract the features of the image and represent them as multidimensional vectors, which, when summed up, will allow you to reconstruct the original image. Having started with a single object in the scene, we will try to move to a set of several objects.

2 Datasets

The task consists of three consecutive steps for each of which a different dataset was created. These datasets are based on the dSprites dataset Figure 1. dSprites is a dataset of 2D shapes procedurally generated from 6 ground truth independent latent factors Table 1




Figure 1: Example of images from dSprites dataset

Table 1: List of features

Feature	Distribution
Shape	square, ellipse, heart
Scale	6 values linearly spaced in $[0.5, 1]$
Orientation	40 values in $[0, 2 \pi]$
Position X	32 values in $[0, 1]$
Position Y	32 values in $[0, 1]$

2.1 Paired-dSprites dataset

The purpose of creating this dataset is to make it possible to easily obtain paired images that differ in a single feature. This is possible due to the fact that in the original dataset the images are arranged in an orderly manner. An example of pairwise images in a dataset can be seen in the Figure 2




Figure 2: Visualizing Paired dSprites dataset elements

2.2 Scene-dsprites dataset

This dataset was created to test the model’s ability to reconstruct a scene from the sum of object vectors. Dataset consists of 2 to 5 non-overlapping figures from dSprites dataset. An example of such images on Figure 3. This version differs from the original multi-dsprites dataset in that the figures do not overlap and have the same color.



Figure 3: Example of collected scenes in the scene-dsprites dataset

2.3 Paired-Scene-dSprites dataset

This dataset combines the capabilities of the first and second dataset. There are two objects in the scene image. One of these objects can change one feature. Figure 4




Figure 4: Example of images from paired-scene-dsprites dataset. From left to right, first scene, pair scene, the object to be changed, his pair, second object

3 Models

The stages of the work involve successive complications of the model. The first stage is that the model will be able to reconstruct the original scene from the sum of features (Figure 5). The second stage - the model will be able to reconstruct the scene from the sum of objects (Figure 6). The third stage - combination of the first two approaches - the model should reconstruct the scene from the sum of objects, which are represented by the sum of features (Figure 7).

A variation autoencoder is chosen as the basic model. The encoder and decoder consist of 4 convolutional and 2 linear layers. The latent representation of one figure in dsprites consists of five 1024-dimensional vectors. One for each feature (Figure 8).

3.1 Paired-dSprites model

This model reconstructs the scene from the sum of feature vectors. The main difficulty is that the vectors representing the features carry information only

about this feature, as well as that the sum of these vectors is reconstructed correctly. Ideally, you want to arrive at a model that can assemble a single object with the desired properties from different attributes of several objects.

In order for the model to learn the matching of a particular latent vector with an appropriate feature, the following training method is proposed. Two images that differ by a single feature are fed to the input of the model, and then both images are passed through the encoder, obtaining a latent representation.

We want the four vectors of latent representation corresponding to the four features to be independent of the changing feature. Therefore it is assumed that one of the vectors encodes the difference between the images. To teach this to the model, we swap these difference vectors between the input features and try to reconstruct the original images crosswise (Figure 9).

3.2 Scene-dSprites model

3.3 Paired-Scene-dSprites model




Figure 5: The summation of image features to obtain a latent representation of the object on the scene.




Figure 6: The summation of image features to obtain a latent representation of the object on the scene.

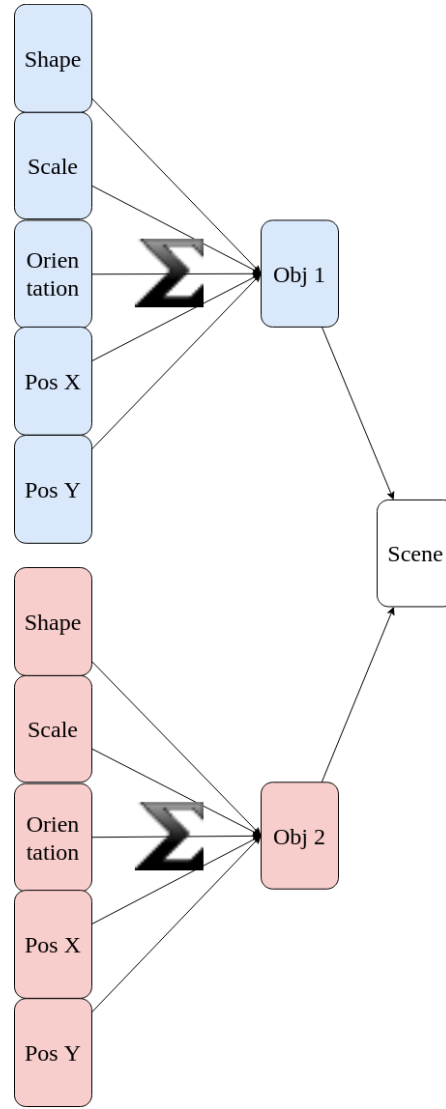


Figure 7: The summation of the latent representations of several objects on the scene.




Figure 8: Diagram of the basic model

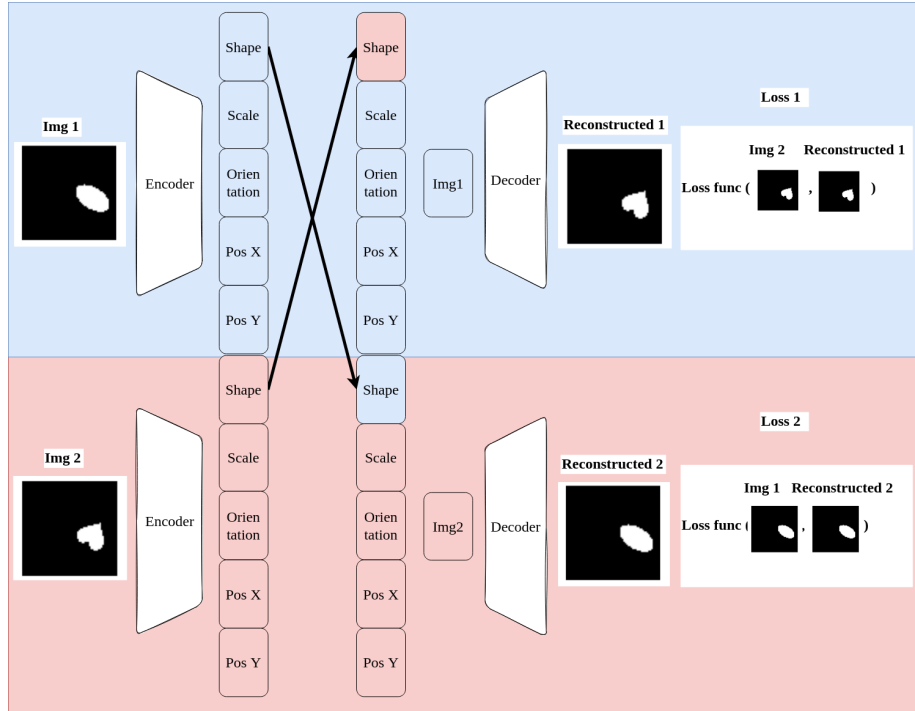


Figure 9: Diagram of the basic model